

CS 5720 Neural Network Deep Learning

ICP-7

CRN	23441
NAME	MANOJ BALA
EMAIL	mxb40210@ucmo.edu
STUDENT_ID	700754021

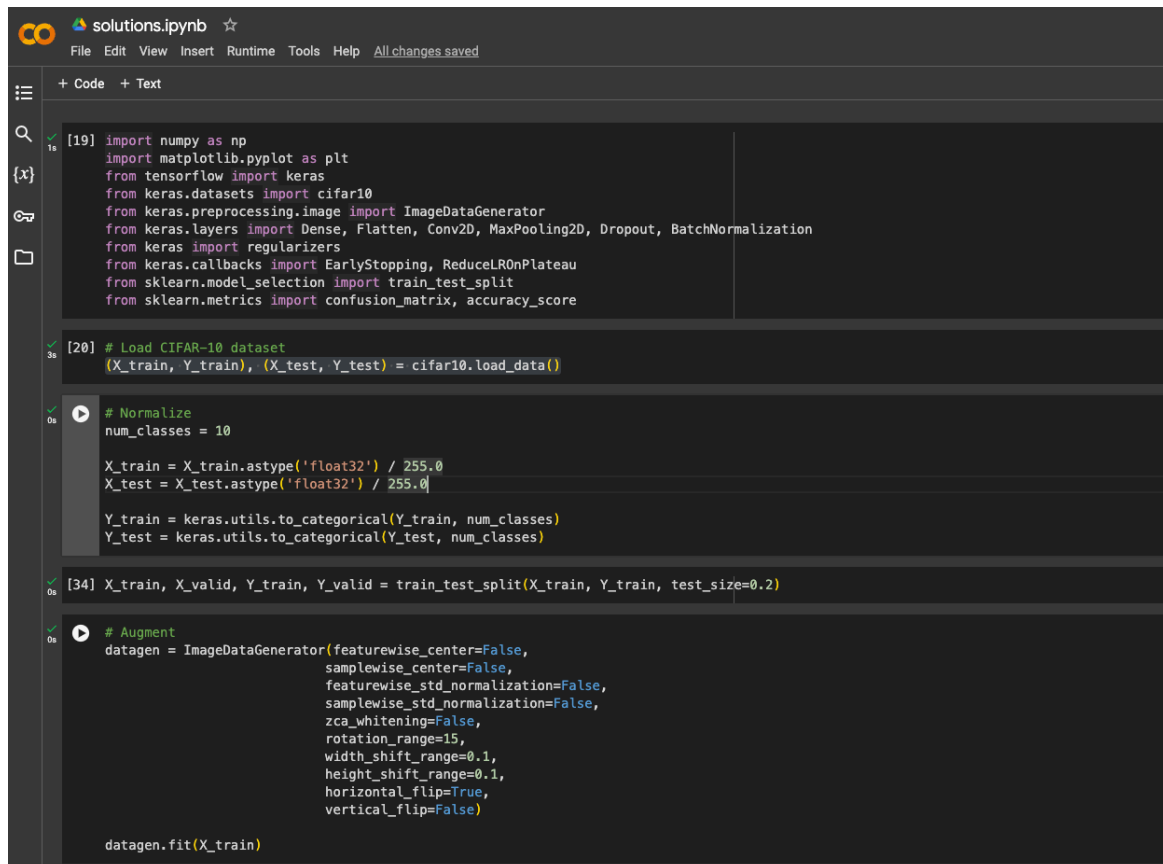
GitHub Repository:

<https://github.com/mxb40210/700754021-NeuralNetworkDeepLearning>

Assignment 7:

<https://github.com/mxb40210/700754021-NeuralNetworkDeepLearning/tree/main/assignments/assignment7>

Screenshots:



```
solutions.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

[19] import numpy as np
import matplotlib.pyplot as plt
from tensorflow import keras
from keras.datasets import cifar10
from keras.preprocessing.image import ImageDataGenerator
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, BatchNormalization
from keras import regularizers
from keras.callbacks import EarlyStopping, ReduceLROnPlateau
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, accuracy_score

[20] # Load CIFAR-10 dataset
(X_train, Y_train), (X_test, Y_test) = cifar10.load_data()

# Normalize
num_classes = 10

X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

Y_train = keras.utils.to_categorical(Y_train, num_classes)
Y_test = keras.utils.to_categorical(Y_test, num_classes)

[34] X_train, X_valid, Y_train, Y_valid = train_test_split(X_train, Y_train, test_size=0.2)

# Augment
datagen = ImageDataGenerator(featurewise_center=False,
                             samplewise_center=False,
                             featurewise_std_normalization=False,
                             samplewise_std_normalization=False,
                             zca_whitening=False,
                             rotation_range=15,
                             width_shift_range=0.1,
                             height_shift_range=0.1,
                             horizontal_flip=True,
                             vertical_flip=False)

datagen.fit(X_train)
```

```
solutions.ipynb
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

# learning_rates = [0.001, 0.0001]
# optimizers = ['adam', 'rms']
# batch_sizes = [32, 64]

learning_rates = [0.0001]
optimizers = ['adam']
batch_sizes = [64]
epochs = 10

# Function to create and train models with different hyperparameters
def create_and_train_model(optimizer='adam', learning_rate=0.001, batch_size=128):
    # Model
    model = keras.Sequential([
        Conv2D(128, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l2(1e-4), input_shape=(32, 32, 3)),
        MaxPooling2D((2, 2)),
        Dropout(0.3),

        Conv2D(256, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l2(1e-4)),
        MaxPooling2D((2, 2)),
        Dropout(0.3),

        Conv2D(512, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l2(1e-4)),
        Conv2D(512, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l2(1e-4)),
        Conv2D(256, (3, 3), activation='relu', padding='same', kernel_regularizer=regularizers.l2(1e-4)),
        MaxPooling2D((2, 2)),
        Dropout(0.3),

        Flatten(),

        Dense(512, activation='relu'),
        Dropout(0.5),
        Dense(256, activation='relu'),
        Dropout(0.5),
        Dense(128, activation='relu'),
        Dropout(0.5),

        Dense(10, activation='softmax')
    ])

    # Compile
    model.compile(optimizer=keras.optimizers.Adam(lr=learning_rate), loss='categorical_crossentropy', metrics=['accuracy'])

    history = model.fit_generator(datagen.flow(X_train, Y_train, batch_size = batch_size),
                                steps_per_epoch = len(X_train) // batch_size,
                                epochs = epochs,
                                validation_data = (X_valid, Y_valid),
                                verbose=1)

    return model, history

# Tuning hyperparameters and saving the histories
models, histories = [], []
for lr in learning_rates:
    for opt in optimizers:
        for bs in batch_sizes:
            model, history = create_and_train_model(opt, learning_rate=lr, batch_size=bs)
            models.append(model)
            histories.append(history)
```

solutions.ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

[36]

WARNING:absl:`lr` is deprecated in Keras optimizer, please use `learning_rate` or use the legacy optimizer, e.g.,tf.keras.optimizers.legacy

Epoch 1/10

<ipython-input-36-4c504b9013ab>:43: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use

history = model.fit_generator(datagen.flow(X_train, Y_train, batch_size = batch_size),

625/625 [=====] - 33s 48ms/step - loss: 2.3075 - accuracy: 0.0997 - val_loss: 2.3030 - val_accuracy: 0.0965

Epoch 2/10

625/625 [=====] - 30s 48ms/step - loss: 2.3029 - accuracy: 0.1015 - val_loss: 2.3027 - val_accuracy: 0.1004

Epoch 3/10

625/625 [=====] - 32s 50ms/step - loss: 2.3028 - accuracy: 0.0983 - val_loss: 2.3026 - val_accuracy: 0.1004

Epoch 4/10

625/625 [=====] - 30s 47ms/step - loss: 2.3029 - accuracy: 0.0975 - val_loss: 2.3029 - val_accuracy: 0.0963

Epoch 5/10

625/625 [=====] - 30s 47ms/step - loss: 2.3028 - accuracy: 0.0976 - val_loss: 2.3027 - val_accuracy: 0.0965

Epoch 6/10

625/625 [=====] - 30s 48ms/step - loss: 2.3027 - accuracy: 0.1007 - val_loss: 2.3027 - val_accuracy: 0.1004

Epoch 7/10

625/625 [=====] - 31s 50ms/step - loss: 2.3028 - accuracy: 0.0996 - val_loss: 2.3027 - val_accuracy: 0.1004

Epoch 8/10

625/625 [=====] - 30s 49ms/step - loss: 2.3028 - accuracy: 0.0990 - val_loss: 2.3027 - val_accuracy: 0.0987

Epoch 9/10

625/625 [=====] - 29s 46ms/step - loss: 2.3027 - accuracy: 0.0979 - val_loss: 2.3028 - val_accuracy: 0.0965

Epoch 10/10

625/625 [=====] - 30s 48ms/step - loss: 2.3027 - accuracy: 0.0979 - val_loss: 2.3027 - val_accuracy: 0.0963

[46]

Get the best model

Extract validation accuracy values from each history object

val_acc_list = [h.history['val_accuracy'][-1] for h in histories]

Find the index of the model with the highest validation accuracy

best_model_idx = np.argmax(val_acc_list)

best_model = models[best_model_idx]

best_history = histories[best_model_idx]

Evaluate the best model on the test set

test_datagen = ImageDataGenerator()

test_datagen.fit(X_test)

loss, acc = best_model.evaluate(test_datagen.flow(X_test, Y_test, batch_size=32))

print('Test Loss: {}, Test Accuracy: {}'.format(loss, acc))

Save the best model for future use

best_model_name = 'best_model.keras'

best_model.save(best_model_name)

print('Model: {} saved!'.format(best_model_name))

313/313 [=====] - 3s 9ms/step - loss: 2.3026 - accuracy: 0.1000

Test Loss: 2.3026161193847656, Test Accuracy: 0.10000000149011612

Model: best_model.keras saved!

[6]

Load the best model

from keras.models import load_model

loaded_model = load_model(best_model_name)

Get predictions on the test set

y_pred = loaded_model.predict(X_test)

y_pred = np.argmax(y_pred, axis=1)

313/313 [=====] - 2s 5ms/step

[39]

Function for Confusion Matrix

def plot_confusion_matrix(cm, class_names):

plt.figure(figsize=(8, 8))

```
solutions.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

[38] # Load the best model
from keras.models import load_model
loaded_model = load_model(best_model_name)

# Get predictions on the test set
y_pred = loaded_model.predict(X_test)
y_pred = np.argmax(y_pred, axis=1)

313/313 [=====] - 2s 5ms/step

# Function for Confusion Matrix
def plot_confusion_matrix(cm, class_names):
    plt.figure(figsize=(8, 8))
    plt.imshow(cm, interpolation='nearest', cmap=plt.cm.Blues)
    plt.title('Confusion Matrix')
    plt.colorbar()
    tick_marks = np.arange(len(class_names))
    plt.xticks(tick_marks, class_names, rotation=45)
    plt.yticks(tick_marks, class_names)

    # Normalize confusion matrix values to display percentages
    cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    for i, j in np.ndenumerate(cm):
        row, col = i
        value = f"{cm[row, col]:.2f}"
        plt.text(col, row, value, ha='center', va='center', fontsize=8, fontweight='bold', color='black')

    plt.grid(False)
    plt.tight_layout()
    plt.show()

# Function for Loss & Accuracy (subplot)
def plot_loss_accuracy_using_subplot():
    plt.figure(figsize=(12, 6))

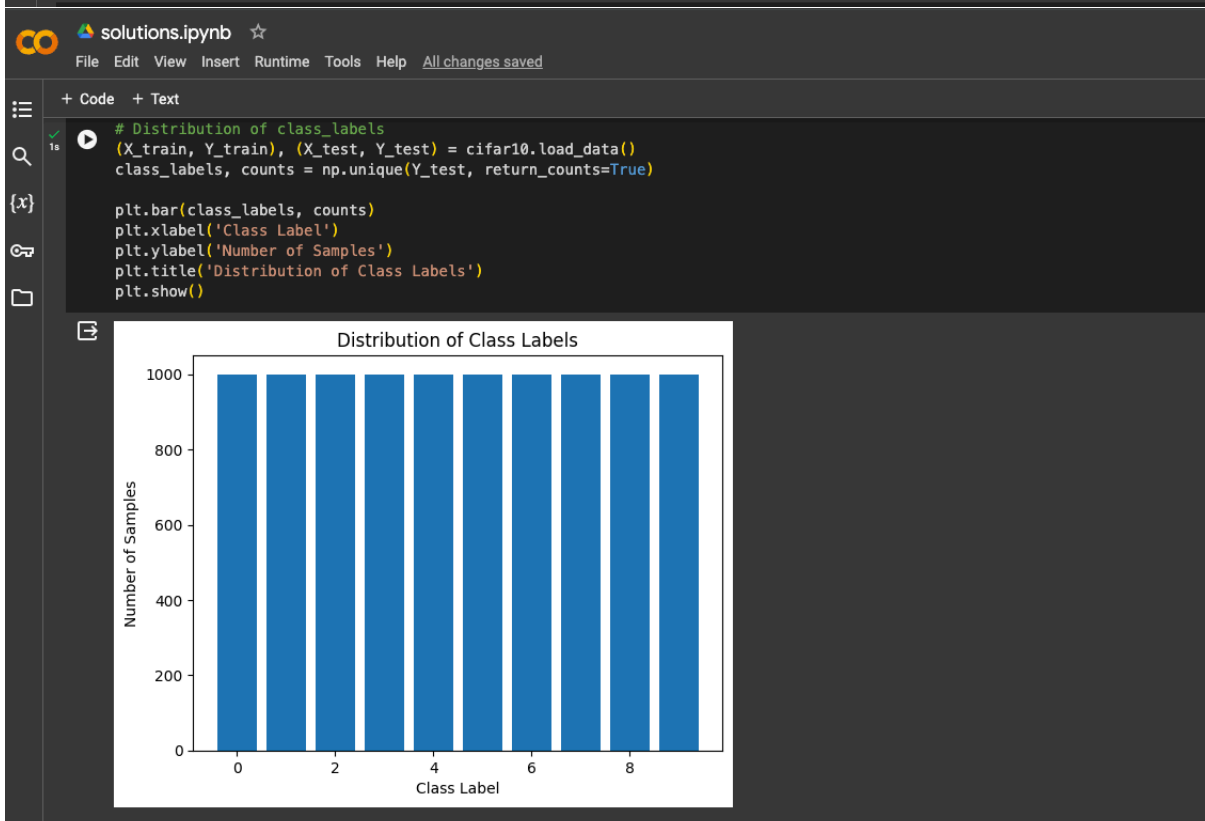
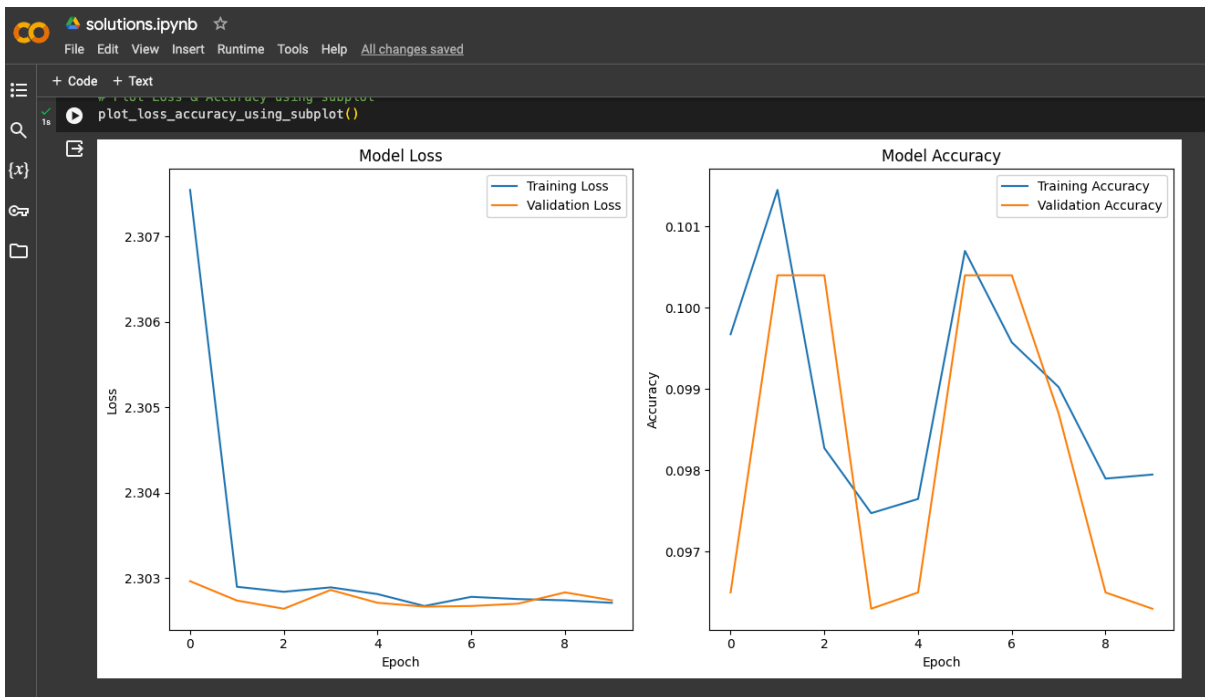
    plt.subplot(2, 1, 1)
    plt.plot(best_history.history['loss'], label='Training Loss')
    plt.plot(best_history.history['val_loss'], label='Validation Loss')
    plt.title('Model Loss')
    plt.ylabel('Loss')
    plt.xlabel('Epoch')
    plt.legend()

    plt.subplot(2, 1, 2)
    plt.plot(best_history.history['accuracy'], label='Training Accuracy')
    plt.plot(best_history.history['val_accuracy'], label='Validation Accuracy')
    plt.title('Model Accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.legend()

    plt.tight_layout()
    plt.show()

[27] # Classes
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

# Confusion Matrix
cm = confusion_matrix(Y_test, y_pred)
```



solutions.ipynb

☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

✓

0s

```
# Assuming you have sample images (X_test), true labels (Y_test), and predicted labels (y_pred)
num_images = 4 # Adjust the number of images

fig, axes = plt.subplots(2, 2, figsize=(3, 3))


for i in range(num_images):
    image = X_test[i]
    true_label = class_labels[Y_test[i]]
    predicted_label = class_labels[y_pred[i]]

    axes[i // 2, i % 2].imshow(image)
    axes[i // 2, i % 2].set_title(f"True: {true_label}\nPredicted: {predicted_label}")
    axes[i // 2, i % 2].axis('off')


fig.suptitle(f"Sample Images with True and Predicted Labels")
plt.tight_layout()
plt.show()
```

Sample Images with True and Predicted Labels


True: [3]
Predicted: 0




True: [8]
Predicted: 0



True: [8]
Predicted: 0



True: [0]
Predicted: 0





+ Code + Text

✓
15s

```
## On MNIST data

from tensorflow import keras
from keras.datasets import mnist
from keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout
from keras.callbacks import EarlyStopping, ReduceLR0nPlateau
import matplotlib.pyplot as plt # Import for plotting

# Load MNIST data
(X_train, Y_train), (X_test, Y_test) = mnist.load_data()

# Reshape to add channel dimension
X_train = X_train.reshape(-1, 28, 28, 1)
X_test = X_test.reshape(-1, 28, 28, 1)

# Normalize pixel values
X_train = X_train.astype('float32') / 255.0
X_test = X_test.astype('float32') / 255.0

# Define learning rate, batch size, and epochs
learning_rate = 0.001
batch_size = 64
epochs = 3

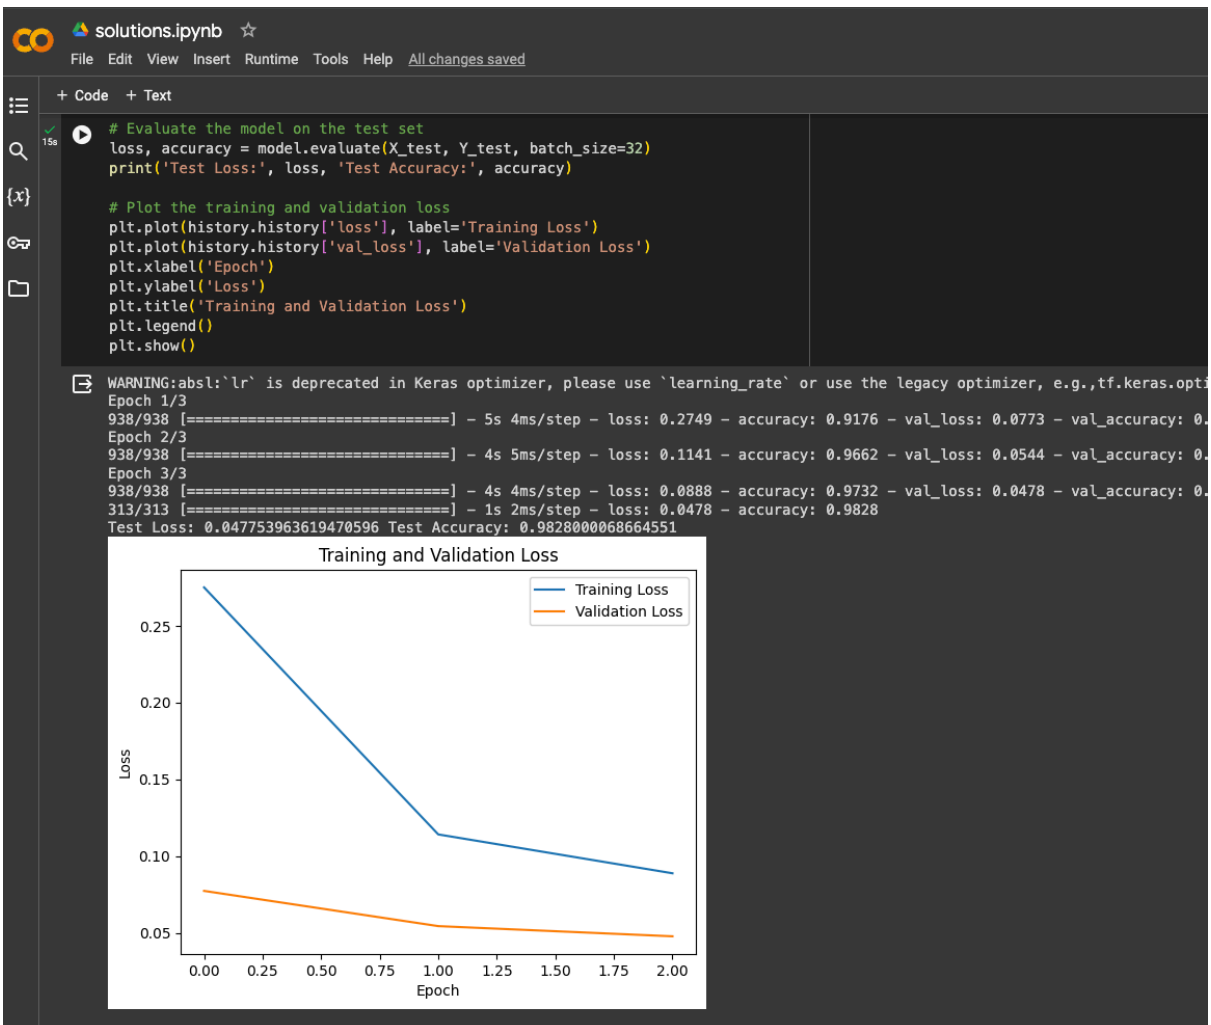
# Build the model
model = keras.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Dropout(0.25),
    Flatten(),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(10, activation='softmax') # 10 units for 10 classes
])

# Compile the model
model.compile(optimizer=keras.optimizers.Adam(lr=learning_rate),
              loss='sparse_categorical_crossentropy', # Categorical crossentropy for MNIST
              metrics=['accuracy'])

# Early stopping and learning rate reduction
early_stopping = EarlyStopping(monitor='val_loss', patience=3)
reduce_lr = ReduceLR0nPlateau(monitor='val_loss', factor=0.5, patience=2, min_lr=0.0001)

# Train the model (without ImageDataGenerator)
history = model.fit(X_train, Y_train,
                    epochs=epochs,
                    validation_data=(X_test, Y_test),
                    batch_size=batch_size,
                    verbose=1,
                    callbacks=[early_stopping, reduce_lr])

# Evaluate the model on the test set
loss, accuracy = model.evaluate(X_test, Y_test, batch_size=32)
print('Test Loss:', loss, 'Test Accuracy:', accuracy)
```





+ Code + Text



```
# Create confusion matrix
cm = confusion_matrix(Y_test, y_pred)

# Function to plot the confusion matrix (optional)
def plot_confusion_matrix(cm, classes, normalize=False, title='Confusion Matrix', cmap=plt.cm.Blues):
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
        print("Normalized confusion matrix")
    else:
        print("Confusion matrix without normalization")
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i in range(len(classes)):
        for j in range(len(classes)):
            plt.text(j, i, format(cm[i, j], fmt),
                     ha="center", va="center",
                     color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True Label')
    plt.xlabel('Predicted Label')
    plt.show()

# Plot confusion matrix (optional)
plot_confusion_matrix(cm, classes=['0', '1', '2', '3', '4', '5', '6', '7', '8', '9'])
```

313/313 [=====] - 1s 2ms/step
Confusion matrix without normalization

